

Report.

5. In 'cqlsh', create a keyspace called 'killrvideo' and switch to that keyspace. Use 'SimpleStrategy' for the replication class with a replication factor of one. Remember the 'use' command switches keyspaces.

NOTE: You can press the tab key within the CREATE KEYSPACE command to have 'cqlsh' autocomplete the replication parameters.

```
cqlsh> CREATE KEYSPACE IF NOT EXISTS KILLERVIDEOS WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
cqlsh> desc keyspaces;
```

killervideos	system_auth	system_distributed	killrvideo
system_schema	system	system_traces	

6. Create a single table called 'videos' with the same structure as shown in table above. 'video_id' is the primary key.

We can switch into the **killervideos** keyspace by doing **use killervideos;** at command prompt.

```
cqlsh:killervideos> CREATE TABLE IF NOT EXISTS videos (
    ... video_id timeuuid PRIMARY KEY,
    ... added_date timestamp,
    ... description text,
    ... title text,
    ... user_id uuid
    ... );
cqlsh:killervideos> desc killervideos;

CREATE KEYSPACE killervideos WITH replication = {'class': 'SimpleStrategy', 'replication_factor': '1'} AND durable_writes = true;

CREATE TABLE killervideos.videos (
    video_id timeuuid PRIMARY KEY,
    added_date timestamp,
    description text,
    title text,
    user_id uuid
) WITH bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND comment = ''
    AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'}
    AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
    AND crc_check_chance = 1.0
    AND dclocal_read_repair_chance = 0.1
    AND default_time_to_live = 0
    AND gc_grace_seconds = 864000
    AND max_index_interval = 2048
    AND memtable_flush_period_in_ms = 0
    AND min_index_interval = 128
    AND read_repair_chance = 0.0
    AND speculative_retry = '99PERCENTILE';

cqlsh:killervideos> █
```

7. Load the newly created table with the 'videos.csv' file using the 'COPY' command.

COPY videos FROM 'videos.csv' WITH HEADER=true;

```
cqlsh:killervideos> COPY videos FROM '/home/pavel_orekhov/labwork/exercise-2/videos.csv' WITH HEADER = true;
Using 1 child processes

Starting copy of killervideos.videos with columns ['video_id', 'added_date', 'description', 'title', 'user_id'].
Processed: 430 rows; Rate: 704 rows/s; Avg. rate: 1031 rows/s
430 rows imported from 1 files in 0.417 seconds (0 skipped).
```

8. Use SELECT to verify the data loaded correctly. Include LIMIT to retrieve only the first 10 rows.

```
cqlsh:killervideos> SELECT * FROM videos LIMIT 10;
```

video_id title	added_date	description	user_id
26461a70-14bd-11e5-ad08-843835b7e3a	2014-05-07 00:00:00.000000+0000	At Comcast we are working on the future of television. Webinar: Building Blocks for the Future of Television	10d5c76c-8767-4db3-8050-e19e
2645e79c-14bd-11e5-a456-843835b7e3a	2011-10-21 00:00:00.000000+0000	DataStax is the developer of DataStax Enterprise, a distributed database. DataStax Cassandra Tutorials - Understanding partitioning and replication in Cassandra	10d5c76c-8767-4db3-8050-e19e
9056808b-ca65-1bfb-9957-3beal48dfdce	2015-03-09 00:00:02.000000+0000	New hire Chip (Chris Hemsworth) learns it's not all about the Empire. Empire Promo - SNL	220077ff-be79-4f20-8603-1b9e
264601a3-14bd-11e5-8c2e-843835b7e3a	2011-12-30 00:00:00.000000+0000	Tyler Hobbs - Flexibility: Python Clients for Apache Cassandra	10d5c76c-8767-4db3-8050-e19e
fe3c4045-6f37-1223-81be-250dc60cfc8	2015-01-16 22:46:44.000000+0000	Saturday Night Live celebrates 40 Years in the Making - Saturday Night Live	539fd1b2-ff34-42c9-80cd-a34e
2e8ecb4f-e92b-139b-8183-4df0e2a817bb	2015-04-24 00:00:41.000000+0000	As new types of data sources emerge from cloud, mobile and social, it's important to optimize and simplify database performance. Webinar: Don't leave your data in the dark - Optimize and simplify database performance	53a8ea04-018b-44c2-a420-c05e
2646123a-14bd-11e5-b9db-843835b7e3a	2012-08-20 00:00:00.000000+0000	C* 2012: Cassandra in Action - Solving Big Data Problems (Eddie Satterly, Splunk)	10d5c76c-8767-4db3-8050-e19e
bd57288-e51c-1ff1-805d-c5f1e49c2c8b	2015-01-19 08:00:00.000000+0000	Join Helena Edelson, Senior Software Engineer at Twitter, as she discusses the challenges of scaling a data-driven application. Webinar Streaming Big Data Analytics with Team Apache Spark & Spark Streaming, Kafka, Cassandra	66b25618-683a-4603-b3c9-caa8
2646278f-14bd-11e5-88ea-843835b7e3a	2012-04-27 00:00:00.000000+0000	NoSQL is addressing some tough challenges that businesses face. Webinar: Top 5 gotchas that prevent NoSQL from meeting business goals	10d5c76c-8767-4db3-8050-e19e
607df86e-2208-18a8-90aa-6d837c659f2f	2015-02-09 00:00:01.000000+0000	Delegation allows you to use the user authentication and authorization of a 3rd party. Delegating User Authentication and Product Subscription to a 3rd Party	723f6f5f-3658-4449-90d0-4391

(10 rows)

9. Use SELECT to COUNT(*) the number of imported rows. It should match the number of rows COPY reported as imported.

```
cqlsh:killervideos> SELECT COUNT(*) FROM videos;
```

count
430

(1 rows)

Warnings :

Aggregation query used without partition key

10. Use SELECT to find a row where the video_id = 6c4cffb9-0dc4-1d59-af24-c960b5fc3652. Next we will explore some other CQL commands that will come in handy, like TRUNCATE in a later exercise, we will show you how to add/remove (non-primary key) columns.

```
cqlsh:killervideos> SELECT * FROM videos WHERE video_id=6c4cffb9-0dc4-1d59-af24-c960b5fc3652;
```

video_id	added_date	description	title	user_id
6c4cffb9-0dc4-1d59-af24-c960b5fc3652	2014-11-06 01:11:50.000000+0000	Speaker: Luke Tillman, Language Evangelist at DataStax. Cat--it seems like the Internet can't get enough cat videos. If you were building an application to let users share and consume cat videos, you'd need a scalable, distributed database. I'll take a look at the data model for KillrVideo, a sample video sharing application similar to YouTube where users can share and consume cat videos. I'll show you how to model the data, how the application drives the data model, and how to shift your thinking. Cassandra Day Denver 2014: A Cassandra Data Model for Serving up Cat Videos		fd3f7889-fc0c-43db-951c-7b77710898bc

(1 rows)

11. Let's remove the data from our table using TRUNCATE .truncate videos;

```
cqlsh:killervideos> TRUNCATE videos;
```

```
cqlsh:killervideos> SELECT * FROM videos;
```

video_id	added_date	description	title	user_id
----------	------------	-------------	-------	---------

(0 rows)

12. Create a second table in the 'killvideo' keyspace called 'videos_by_title_year' with the structure shown in above table.

```
cqlsh:killervideos> CREATE TABLE IF NOT EXISTS videos_by_title_year (
    ... video_id timeuuid,
    ... added_year int,
    ... added_date timestamp,
    ... description text,
    ... title text,
    ... user_id uuid,
    ... PRIMARY KEY((title, added_year))
    ... );
cqlsh:killervideos> COPY videos_by_title_year FROM '/home/pavel_orekhov/labwork/exercise-3/videos_by_title_year.csv' WITH HEADER=true;
Using 1 child processes

Starting copy of killervideos.videos_by_title_year with columns ['title', 'added_year', 'added_date', 'description', 'user_id', 'video_id'].
Processed: 430 rows; Rate: 678 rows/s; Avg. rate: 1025 rows/s
430 rows imported from 1 files in 0.419 seconds (0 skipped).
```

An extra set of parentheses is necessary, to combine **title** and **added_year** into 1 column.

13. Load the data from the 'videos_by_title_year.csv' file using the 'COPY' command.

14. COPY videos_by_title_year FROM 'videos_by_title_year.csv' WITH HEADER=true;

```
cqlsh:killervideos> COPY videos_by_title_year FROM '/home/pavel_orekhov/labwork/exercise-3/videos_by_title_year.csv' WITH HEADER=true;
Using 1 child processes

Starting copy of killervideos.videos_by_title_year with columns ['title', 'added_year', 'added_date', 'description', 'user_id', 'video_id'].
Processed: 430 rows; Rate: 639 rows/s; Avg. rate: 977 rows/s
430 rows imported from 1 files in 0.440 seconds (0 skipped).
cqlsh:killervideos>
```

15. Try running queries on the 'videos_by_title_year' table to query on a specific 'title' and 'added_year'.

```
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: added_year must be restricted as other parts are"
cqlsh:killervideos> select * from videos_by_title_year where added_year = 1999;
InvalidRequest: code=2200 [Invalid query] message="Partition key parts: title must be restricted as other parts are"
cqlsh:killervideos>
```

16. What error does Cassandra return when you try to query on just title or just year? Why?

It returns the error message I posted above. Because they have been combined into 1 compound partition key, hence the hash value for the partition was computed on this paired object, and not on the individual columns, hence, if you give it just one part of the paired object, it won't be able to retrieve the hash value of the pair.

17. Create a table with the columns above to facilitate querying for videos by tag within a given year range returning the results in descending order by year.

18. We wrote most of the CREATE TABLE for you. Fill in the PRIMARY KEY and CLUSTERING ORDER BY.

```
CREATE TABLE videos_by_tag_year (
    tag text,
    added_year int,
    video_id timeuuid,
    added_date timestamp,
    description text,
    title text,
```

```
    user_id uuid,
    PRIMARY KEY ( )
) WITH CLUSTERING ORDER BY ( );
```

```
cqlsh:killervideos> CREATE TABLE IF NOT EXISTS videos_by_tag_year(
... tag text,
... added_year int,
... video_id timeuuid,
... added_date timestamp,
... description text,
... title text,
... user_id uuid,
... PRIMARY KEY(tag, added_year, video_id)
... ) WITH CLUSTERING ORDER BY(added_year DESC);
```

I added **video_id** into the primary key for simplicity, when copying data, otherwise, it is not kept in the right order.

19. Load the data from the 'videos_by_tag_year.csv' file in the provided 'exercise=4' directory using the COPY command.
COPY videos_by_tag_year FROM 'videos_by_tag_year.csv' WITH HEADER=true;

```
cqlsh:killervideos> COPY videos_by_tag_year FROM '/home/pavel_orekhov/labwork/exercise-4/videos_by_tag_year.csv' WITH HEADER=true;
Using 1 child processes

Starting copy of killervideos.videos_by_tag_year with columns ['tag', 'added_year', 'video_id', 'added_date', 'description', 'title', 'user_id'].
Processed: 797 rows; Rate: 1217 rows/s; Avg. rate: 1853 rows/s
797 rows imported from 1 files in 0.430 seconds (0 skipped).
```

20. Check the number of rows in the 'videos_by_tag_year' table.

NOTE: The number of rows should match the number of rows imported by the COPY command.

```
cqlsh:killervideos> SELECT COUNT(*) FROM videos_by_tag_year;

count
-----
797

(1 rows)

Warnings :
Aggregation query used without partition key
```

21. Try running queries on the 'videos_by_tag_year' table to query on a specific tag and added year.

Example queries:

Tag	Added_year
trailer	2015
cql	2014
spark	2014

```
cqlsh:killervideos> SELECT COUNT(*) FROM videos_by_tag_year WHERE tag = 'trailer' AND added_year = 2015;

count
-----
15

(1 rows)
```

```
cqlsh:killervideos> SELECT COUNT(*) FROM videos_by_tag_year WHERE tag = 'cql' AND added_year = 2014;

count
-----
1

(1 rows)
```

```
cqlsh:killervideos> SELECT COUNT(*) FROM videos_by_tag_year WHERE tag = 'spark' AND added_year = 2014;

count
-----
1
```

22. Try querying for all videos with tag "cql" added before the year 2015. Notice you can do range queries on clustering columns.

```
cqlsh:killervideos> SELECT title FROM videos_by_tag_year WHERE tag = 'cql' AND added_year < 2015;

title
-----
The Last Pickle: Lesser Known Features of Cassandra 2.0 and 2.1

(1 rows)
```

23. Try querying for all videos added before 2015. The query will fail. What error message does cqlsh report? Why did the query fail whereas the previous query worked?

```
cqlsh:killervideos> SELECT title FROM videos_by_tag_year WHERE added_year < 2015;
InvalidRequest: code=2200 [Invalid query] message="Cannot execute this query as it might involve data filtering and thus may have unpredictable performance. If you want to execute this query despite the performance unpredictability, use ALLOW FILTERING"
```

Cassandra gives this message as a form of warning, telling us that the performance of our query will be bad and utilizing a lot of computing resources. Indeed, it will have to go through the entire table, filtering out the **added_year** values that we do not want, which is suboptimal in the case when there are a lot of such unwanted values.

The previous query worked, because we filter on the field that our table is partitioned by, which is supposed to return only a small portion of data. It did not complain about the **added_year** range filter, because our rows are ordered by the **added_year** column within each partition.

25. Run the TRUNCATE command to erase the data from the 'videos' table.

```
cqlsh:killervideos> TRUNCATE videos;
```

26. Alter the 'videos' table to add a 'tags' column.

```
cqlsh:killervideos> ALTER TABLE videos ADD tags text;
cqlsh:killervideos>
```

27. Load the data from the 'videos.csv' file using the COPY command. COPY videos FROM 'videos.csv' WITH HEADER=true; Remember, we do not need to create the user defined type called 'video_encoding' because we did so in the previous exercise. However, take a look at the code below as a refresher. Do not run it again or you will get an error!

```
CREATE TYPE video_encoding (  
    bit_rates SET<TEXT>,  
    encoding TEXT,  
    height INT,  
    width INT,  
);
```

```
cqlsh:killervideos> COPY videos FROM '/home/pavel_orekhov/labwork/exercise-5/videos.csv' with HEADER = true;  
Using 1 child processes  
  
Starting copy of killervideos.videos with columns ['video_id', 'added_date', 'description', 'tags', 'title', 'user_id'].  
Processed: 430 rows; Rate: 759 rows/s; Avg. rate: 1112 rows/s  
430 rows imported from 1 files in 0.387 seconds (0 skipped).  
cqlsh:killervideos>
```

28. Alter your table to add an 'encoding' column of the 'video_encoding' type.

```
cqlsh:killervideos> ALTER TABLE videos ADD encoding frozen<video_encoding>;  
cqlsh:killervideos>
```

29. Load the data from the 'videos_encoding.csv' file using the COPY command.
COPY videos (video_id, encoding) FROM 'videos_encoding.csv' WITH HEADER=true;

```
cqlsh:killervideos> COPY videos (video_id, encoding) FROM '/home/pavel_orekhov/labwork/exercise-5/videos_encoding.csv' WITH HEADER = true;  
Using 1 child processes  
  
Starting copy of killervideos.videos with columns ['video_id', 'encoding'].  
Processed: 430 rows; Rate: 376 rows/s; Avg. rate: 628 rows/s  
430 rows imported from 1 files in 0.685 seconds (0 skipped).
```

30. Run a query to retrieve the first 10 rows of the 'videos' table. Notice the altered table contains data for the new 'tags' and 'encoding' column.

```
SELECT * FROM videos LIMIT 10;
```

video_id	added_date	description	encoding	tags	title
1	2012-10-10 10:10:10	an example of our working on the future of television. Though not television are suggesting more reality than ever thanks to the cloud based TV platform ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
2	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
3	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
4	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
5	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
6	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
7	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
8	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
9	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }
10	2012-10-10 10:10:10	the future of television is ...	{'encoding': '1080p', 'height': 1080, 'width': 1920, 'bit_rates': ['10000 Kbps', '4000 Kbps', '1000 Kbps'] }	{'facebook', 'twitter'}	'Facebook', 'Twitter' }

It does indeed.